

ACRONYMS

AOP:	Aspect-Oriented Programming
ALU:	Arithmetic and Logic Unit
API:	Application Programming Interface
ATM:	Asynchronous Transfer Mode
B/S:	Browser-Server
C/S:	Client-Server
CS:	Computer Science
DBMS:	Data Base Management System
FPU:	Float Point Unit
I/O:	Input and Output
ISA:	Instruction Set Architecture
ISO:	International Organization for Standardization
ISP:	Internet Service Providers
KA:	Knowledge Area
LAN:	Local Area Network
MUX:	Multiplexer
NIC:	Network Interface Card
OOP:	Object-Oriented Programming
OS:	Operating Systems
PDA:	Personal Digital Assistant
PDC:	Parallel and Distributed Computing
PL:	Programming Languages
PPP:	Point-to-Point Protocol
RFID:	Radio Frequency Identification
SCSI:	Small Computer System Interface
SQL:	Structured Query Language
VPN:	Virtual Private Network
WAN:	Wide Area Network

Introduction

The Computing Foundations Knowledge Area (KA) is concerned with the computer science foundations that support the design and construction of software products. Moreover, it also includes knowledge about the transformation of a (software) design into a (software) implementation, the tools used during this process, and the various software development methods.

It is generally accepted that software engineering builds on top of computer science. For example, “Software Engineering 2004: Curriculum Guidelines for Undergraduate Degree Programs in Software Engineering” clearly states, “One particularly important aspect is that *software engineering builds on computer science* and mathematics” (italics added).

While some may not feel very comfortable with the claim that software engineering builds on computer science (Frezza) [2], not many will deny the important role computer science plays in the development of software engineering both as a discipline and as a body of knowledge. However, the importance of computer science to software engineering is not sufficiently acknowledged within the software industry; thus, this KA guideline is being written.

Although the term “Computing Foundations” is not precisely defined and may have different meanings or connotations to different people, most would agree that the scope encompassed by the term includes at least the development and operational environment in which software evolves and executes. Because no software can exist in a vacuum or run without a computer, the core of such an environment is the computer and its various components. To some extent, knowledge about the computer and its underlying principles of hardware and software serves as a framework on which software engineering is anchored. Thus, all software engineers must have good understanding of the Computing Foundations Knowledge Area.

The majority of topics discussed in the Computing Foundations KA are also topics of discussion in basic courses given in computer science undergraduate and graduate programs. Such courses include programming, data structure, algorithms, computer organization, operating systems, compiler, databases, networking, distributed systems, and so forth. Thus, when breaking down topics, it can be tempting to decompose the Computing Foundations KA according to these often-found divisions in relevant courses.

However, a pure course-based division of topics suffers serious drawbacks. For one, not all courses in computer science are related or equally important to software engineering. Thus, some topics that would otherwise be covered in a computer science course are not covered in this guideline. For example, computer graphics—while an important course in a computer science degree program—is not included in this guideline.

66 Second, some topics discussed in this guideline do
67 not exist as standalone courses in undergraduate or
68 graduate computer science programs.
69 Consequently, such topics may not be adequately
70 covered in a pure course-based breakdown. For
71 example, abstraction is a topic mentioned/covered
72 in several different computer science courses; it is
73 unclear which course abstraction should belong to
74 in a course-based breakdown of topics.

75 For the aforementioned reasons, we do not use a
76 course-based division; instead, we break the
77 Computing Foundations KA into seventeen different
78 topics. A topic's direct usefulness to software
79 engineers is the criterion we used for selecting
80 topics for inclusion in this KA. (Figure 1 depicts the
81 breakdown of topics for the Computing Foundations
82 KA.) The advantage of this topic-based breakdown is
83 its foundation on the belief that Computing
84 Foundations— if it is to be grasped firmly—must be
85 considered as a collection of logically connected
86 topics undergirding software engineering in general
87 and software construction in particular.

88 The Computing Foundations KA is related closely to
89 the Software Design, Software Construction,
90 Software Testing, Software Maintenance, Software
91 Quality, and the Mathematical Foundations KAs.

92 BREAKDOWN OF TOPICS FOR COMPUTING 93 FOUNDATIONS

94 The breakdown of topics for the Computing
95 Foundations KA is shown in Figure 1.

96 1. Problem Solving Techniques 97 [3*c3, s4.1, s4.2, 4*c5]

98 The concepts, notions, and terminology introduced
99 here form an underlying basis for understanding the
100 role and scope of problem solving techniques.

101 1.1 Definition of Problem Solving

102 Problem solving refers to the thinking and activities
103 conducted to answer or derive a solution to a
104 problem. There are many ways to approach a
105 problem, and each way employs different tools and
106 uses different processes. These different ways of
107 approaching problems gradually expand and define
108 themselves and finally give rise to different
109 disciplines. For example, software engineering
110 focuses on solving problems using computers and
111 software.

112 While different problems warrant different solutions
113 and may require different tools and processes to be
114 used, the methodology and techniques used in
115 solving problems do follow some general guidelines
116 and can often be generalized as problem solving
117 techniques. For example, a general guideline for
118 solving a generic engineering problem is to use the
119 three-step process given below.

- 120 ● Formulate the real problem
- 121 ● Analyze the problem
- 122 ● Design a solution search strategy

123 1.2 Formulating the Real Problem

124 Gerard Voland writes, "It is important to recognize
125 that a specific problem should be formulated if one
126 is to develop a specific solution" [3*p88]. To solve a
127 problem, one must first define the problem. In
128 software jargon, this definition is called the problem
129 statement, which explicitly specifies both what the
130 problem and the desired outcome are.

131 Although there is no specific way of stating a
132 problem, in general a problem should be expressed
133 in such a way as to facilitate the development of
134 solutions for solving it. Some general techniques to
135 help one formulate the real problem include
136 statement-restatement, determining the source and
137 the cause, revising the statement, analyzing present
138 and desired state, and using the fresh eye approach.



Figure 1. Breakdown of topics for the Computing Foundations KA

1.3 Analyze the Problem

Once the problem statement is available, the next step is to analyze the problem statement or situation to help structure our search for a solution. Four types of analysis include *situation analysis*, in which the most urgent or critical aspects of a situation are identified first; *problem analysis*, in which the cause of the problem must be determined; *decision analysis*, in which the action(s) needed to correct the problem or eliminate its cause must be determined; and *potential problem analysis*, in which the action(s) needed to prevent any reoccurrences of the problem or the development of new problems must be determined.

1.4 Design a Solution Search Strategy

Once the problem analysis is complete, we can focus on structuring a search strategy to find the solution. Voland writes, "In most problem solving situations, one should not seek to generate and evaluate every possible solution to a problem" [3*p117]. In order to find the "best" solution, we need to eliminate paths that do not lead to viable solutions, design tasks in a way that provides the most guidance in searching for a solution, and use various attributes of the final solution state to guide our choices in the problem solving process.

1.5 Problem Solving using Programs

The uniqueness of computer software gives problem solving a flavor that is distinct from general engineering problem solving. For example, the techniques employed in solving a programming problem are somewhat different from those used in solving a general engineering problem: to solve a problem using computers, we must answer the following questions.

- How do we figure out what to tell the computer to do?
- How do we convert the problem statement into an algorithm?
- How do we convert the algorithm into machine instructions?

The first task in solving a problem using a computer is to determine what to tell the computer to do. There may be many ways to tell the story, but all should take the perspective of a computer such that the problem can eventually be solved by the computer. In general, a problem should be expressed in such a way as to facilitate the development of algorithms and data structures for solving it.

The result of the first task is a problem statement. The next step is to convert the problem statement into algorithms that solve the problem. Once an algorithm is found, the final step converts the algorithm into machine instructions that form the

197 final solution: software that solves the problem.

198 Abstractly speaking, problem solving using a
199 computer can be considered as a process of
200 problem transformation—in other words, the step-
201 by-step transformation of a problem statement into
202 a problem solution. To the discipline of software
203 engineering, the ultimate objective of problem
204 solving is to transform a problem expressed in
205 natural language into electrons running around a
206 circuit. In general, this transformation can be broken
207 into three phases:

- 208 ● Systematic derivation of algorithms from the
209 problem statement.
- 210 ● Systematic application of algorithms to the
211 problem.
- 212 ● Transformation of algorithms to program code.

213 The conversion of a problem statement into
214 algorithms and algorithms into program codes
215 usually follows a so-called “stepwise refinement”
216 (a.k.a. systematic decomposition) in which we start
217 with a problem statement, rewrite it as a task, and
218 recursively decompose the task into a few simpler
219 subtasks until the task is so simple that solutions to
220 it are straightforward. There are three basic ways of
221 decomposing: sequential, conditional, and iterative.

222 2. Abstraction

223 [4*s5.2-5.4]

224 Abstraction is an indispensable technique associated
225 with problem solving. It refers to both the process
226 and result of generalization by reducing the
227 information of a concept, a problem, or an
228 observable phenomenon so that one can focus on
229 the “big picture.” One of the most important skills in
230 any engineering undertaking is framing the levels of
231 abstraction appropriately.

232 “Through abstraction,” according to Volland, “we
233 view the problem and its possible solution paths
234 from a higher level of conceptual understanding. As
235 a result, we may become better prepared to
236 recognize possible relationships between different
237 aspects of the problem and thereby generate more
238 creative design solutions” [3*p232]. This is
239 particularly true in computer science in general
240 (such as hardware vs. software) and in software
241 engineering in particular (software requirement vs.
242 software construction, and so forth).

243 2.1 Levels of Abstraction

244 When abstracting, we concentrate on one “level” of the
245 big picture at a time with confidence that we can then
246 connect effectively with levels above and below.
247 Although we focus on one level, abstraction does not
248 mean knowing nothing about the neighboring levels.
249 Abstraction levels do not necessarily correspond to
250 discrete components in reality or in the problem
251 domain, but to well-defined standard interfaces such as
252 programming APIs. The advantages that standard
253 interfaces provide include *portability*, *easier*
254 *software/hardware integration* and *wider usage*.

255 Usually, the definition or divisions of abstraction levels
256 are arbitrary to some extent; there are often other ways
257 to draw the lines.

258 2.2 Encapsulation

259 Encapsulation is a mechanism used to implement
260 abstraction. When we are dealing with one level of
261 abstraction, the information concerning the levels
262 below and above that level is encapsulated. This
263 information can be the concept, problem, or
264 observable phenomenon; or it may be the
265 permissible operations on these relevant entities.
266 Encapsulation usually comes with some degree of
267 information hiding in which some or all of the
268 underlying details are hidden from the level above
269 the interface provided by the abstraction. To an
270 object, information hiding means we don't need to
271 know the details of how the object is represented or
272 how the operations on those objects are
273 implemented.

274 2.3 Hierarchy

275 When we use abstraction in our problem
276 formulation and solution, we may use different
277 abstractions at different times—in other words, we
278 work on different levels of abstraction as the
279 situation calls. Most of the time, these different
280 levels of abstraction are organized in a hierarchy.
281 This hierarchical structure resembles the
282 organizational structures in human society and thus
283 facilitates easier understanding by humans. There
284 are many ways to structure a particular hierarchy
285 and the criteria used in determining the specific
286 content of each layer in the hierarchy varies
287 depending on the individuals performing the work.

Most of the time, a hierarchy of abstraction levels is sequential, which means that each layer has one and only one predecessor (lower) layer and one and only one successor (upper) layer—except the upmost layer (which has no successor) and the bottommost layer (which has no predecessor). Sometimes, the hierarchy is organized in a tree-like structure, which means each layer can have more than one predecessor layer but only one successor layer. Occasionally, a hierarchy can have a many-to-many structure, in which each layer can have multiple predecessors and successors. At no time, shall there be any loop in a hierarchy.

A hierarchy often forms naturally in task decomposition. Often, a task analysis can be decomposed in a hierarchical fashion, starting with the larger tasks and goals of the organization and breaking each of them down into smaller subtasks that can again be further subdivided. This continuous division of tasks into smaller ones would produce a hierarchical structure of tasks-subtasks.

3. Programming Fundamentals [4*c6-19]

Programming (often called coding) is composed of the methodologies or activities for creating computer programs that perform a desired function. It is an indispensable part in software construction. In general, programming can be considered as the process of designing, writing, testing, debugging, and maintaining the source code. This source code is written in some programming language.

The process of writing source code often requires expertise in many different subject areas—including knowledge of the application domain, appropriate data structures, specialized algorithms, various language constructs, good programming techniques, and software engineering.

3.1 The Programming Process

Programming involves design, writing, testing, debugging, and maintenance. *Design* is the conception or invention of a scheme for turning a customer requirement for computer software into operational software. It is the activity that links application requirement to coding and debugging.

Writing is the actual coding of the design in an appropriate programming language. *Testing* is the activity to verify that the code one writes actually does what it is supposed to do. *Debugging* is the activity to find and fix bugs (errors) in the source code (or design). *Maintenance* is the activity to update, correct, and enhance existing programs. Each of these activities is a huge topic and often warrants the explanation of an entire Knowledge Area in the SWEBOK Guide and many books.

3.2 Programming Paradigms

Programming is highly creative and thus somewhat personal. Different people often write different programs for the same requirements. This diversity of programming causes much difficulty in the construction and maintenance of large complex software. Various programming paradigms have been developed over the years to put some standardization into this highly creative and personal activity. When one programs, he or she can use one of many programming paradigms to write the code. The major types of programming paradigms are discussed below.

Unstructured Programming: In unstructured programming, a programmer follows his/her hunch to write the code in whatever way he/she likes as long as the function is operational. Often, the practice is to write code to fulfill a specific utility without regard to anything else. Programs written this way exhibit no particular structure—thus the name “unstructured programming.” Unstructured programming is also sometimes called ad hoc programming.

Structured/Procedural Programming: Programs are structured as procedures (or functions) with each procedure performing a specific task. Standard interfaces exist between procedures to facilitate correct and smooth calling operations of the programs. Under structured programming, programmers often follow established protocols and rules of thumb when writing code. These protocols and rules can be numerous and cover almost the entire scope of programming—ranging from the simplest issue (such as how to name variables, functions, procedures, and so forth) to more complex issues (such as how to structure an interface, how to use exceptions, and so forth).

380 *Object-Oriented Programming:* While procedural
381 programming organizes programs around
382 procedures, object-oriented programming (OOP)
383 organize a program around so-called objects, which
384 are abstract data structures that combine both data
385 and methods used to access or manipulate the data.
386 The primary features of OOP are that objects
387 representing various abstract and concrete entities
388 are created and these objects interact with each
389 other to collectively fulfill the desired functions.

390 *Aspect-Oriented Programming:* Aspect-oriented
391 programming (AOP) is a programming paradigm
392 that is built on top of OOP. AOP aims to isolate
393 secondary or supporting functions from the main
394 program's business logic by focusing on the so-
395 called cross sections (concerns) of the objects. The
396 primary motivation for AOP is to resolve the so-
397 called object tangling and scattering associated with
398 OOP, in which the interactions among objects
399 become very complex. The essence of AOP is the
400 greatly emphasized separation of concerns, which
401 separates non-core functional concerns or logic into
402 various aspects.

403 3.3 Defensive Programming

404 In the construction of large complex software, many
405 people work as a team with each person being
406 responsible for part of the software. Because
407 different people have different styles in writing
408 programs, it is imperative that people write code in
409 such a way as to avoid conflicts or problems with
410 code written by other people. This thinking leads to
411 the emergence of the so-called defensive
412 programming.

413 Abstractly speaking, defensive programming is more
414 a programming attitude than a style. It takes
415 responsibility for protecting one's own code even
416 when someone else is not doing his/her work
417 properly. The main idea is that if someone does
418 something bad—such as provide a bad input—the
419 code does not break. In a practical sense, defensive
420 programming can be seen as a collection of
421 programming techniques that, when used, produce
422 more robust programs. Techniques for defensive
423 programming include assertions, error-handling,
424 exception, and barricades. Defensive programming
425 often produces higher quality code.

426 4. Programming Language Basics 427 [6*c6]

428 Using computers to solve problems involves
429 programming—writing and organizing instructions
430 telling the computer what to do at each step.
431 Programs must be written in some programming
432 language with which and through which we
433 describe necessary computations. In other words,
434 we use the facilities provided by a programming
435 language to describe problems, develop algorithms,
436 and reason about problem solutions. To write any
437 program, one must understand at least one
438 programming language.

439 4.1 Programming Language Overview

440 A programming language is designed to express
441 computations that can be performed by a computer. In
442 a practical sense, a programming language is a notation
443 for writing programs and thus should be able to express
444 most data structures and algorithms. Some, but not all,
445 people restrict the term "programming language" to
446 those languages that can express all possible
447 algorithms.

448 Not all languages have the same importance and
449 popularity. The most popular ones are often defined by
450 a specification document established by a well-known
451 and respected organization. For example, the C
452 programming language is specified by an ISO standard
453 named *ISO/IEC 9899*. Other languages, such as *Perl*
454 and *Python*, do not enjoy such treatment and often have
455 a dominant *implementation* that is used as a reference.

456 4.2 Syntax and Semantics of Programming Languages

457 Just like natural languages, many programming
458 languages have some form of written specification of
459 their syntax (form) and semantics (meaning). Such
460 specifications include, for example, specific
461 requirements for the definition of variables and
462 constants (in other words, declaration and types) and
463 format requirements for the instructions themselves.

464 In general, a programming language supports such
465 constructs as variables, data types, constants, literals,
466 assignment statements, control statements, procedures,
467 functions, and comments. The syntax and semantics of
468 each construct must be clearly specified.

469 4.3 Low-Level Programming Languages

470 Programming language can be classified into two
471 classes: low-level languages and high-level
472 languages. Low-level languages can be understood
473 by a computer with no or minimal assistance and
474 typically include machine languages and assembly
475 languages. A machine language uses ones and zeros
476 to represent instructions and variables, and is
477 directly understandable by a computer. An assembly
478 language contains the same instructions as a
479 machine language but the instructions and variables
480 have symbolic names that are easier for humans to
481 remember.

482 Assembly languages cannot be directly understood
483 by a computer and must be translated into a
484 machine language by a utility program called an
485 *assembler*. There often exists a correspondence
486 between the instructions of an assembly language
487 and a machine language, and the translation from
488 assembly code to machine code is straightforward.
489 For example, “add r1, r2, r3” is an assembly
490 instruction for adding the content of register r2 and
491 r3 and storing the sum into register r1. This
492 instruction can be easily translated into machine
493 code “0001001010000011” (assume the operation
494 code for addition is 0001.)

495 One common trait shared by these two types of
496 language is their close association with the specifics
497 of a type of computer or instruction set architecture
498 (ISA).

499 4.4 High-Level Programming Languages

500 A high-level programming language has a strong
501 abstraction from the details of the computer’s ISA. In
502 comparison to low-level programming languages, it
503 often uses natural-language elements and is thus much
504 easier for humans to understand. Such languages allow
505 symbolic naming of variables, provide expressiveness,
506 and enable abstraction of the underlying hardware. For
507 example, while each microprocessor has its own ISA,
508 code written in a high-level programming language is
509 usually portable between many different hardware
510 platforms. For this reason, most programmers use
511 and most software are written in high-level
512 programming languages. Examples of high-level
513 programming languages include C, C++, C#, and
514 Java,

515 4.5 Declarative vs. Imperative Programming 516 Languages

517 Most programming languages (high-level or low-level)
518 allow programmers to specify the individual
519 instructions that a computer is to execute. Such
520 programming languages are called imperative
521 programming languages because one has to specify
522 every step clearly to the computer. But some
523 programming languages allow programmers to only
524 describe the function to be performed without
525 specifying the exact instruction sequences to be
526 executed. Such programming languages are called
527 declarative programming languages. Declarative
528 languages are high-level languages. The actual
529 implementation of the computation written in such a
530 language is hidden from the programmers and thus is
531 not a concern for them.

532 The key point to note is that declarative programming
533 only describes *what* the program should accomplish
534 without describing *how* to accomplish it. For this
535 reason, many people believe declarative programming
536 facilitates easier software development. Declarative
537 programming languages include Lisp and Prolog, while
538 imperative programming languages include C, C++,
539 and JAVA.

540 5. Debugging Tools and Techniques 541 [4*c23]

542 Once a program is coded, the next step is
543 debugging, which is a methodical process of finding
544 and reducing the number of bugs or errors in a
545 program. The purpose of debugging is to find out
546 why a program doesn’t work. Except for very simple
547 programs, debugging is always necessary.

548 5.1 Types of Errors

549 Usually, a program does not work because it
550 contains bugs or errors that can be either syntactic
551 errors, logical errors, or data errors.

552 *Syntax errors* are typing errors that result in illegal
553 operations. This type of error only applies to high-
554 level programming languages and not to machine
555 languages. In a machine language, any bit pattern
556 corresponds to some legal instruction. In high-level
557 programming languages, syntax errors are often
558 caught during the compilation or translation from
559 the high-level language into machine code. For
560 example, in the C/C++ programming language, the
561 statement “123=constant;” contains a syntax error
562 that will be caught by the compiler during
563 compilation.

564 *Logic errors* are semantic errors that result in
565 incorrect computations or program behaviors. Your
566 program is legal, but wrong! So the results do not
567 match the problem statement or user expectations.
568 For example, in the C/C++ programming language,
569 the inline function “int f(int x) {return f(x-1);}” for
570 computing factorial x! is legal but logically incorrect.
571 This type of error cannot be caught by a compiler
572 during compilation and is often discovered through
573 tracing the execution of the program.

574 *Data errors* are input errors that result either in
575 input data that is different from what the program
576 expects or in the processing of wrong data. This
577 type of error can be discovered by testing the
578 program with a wide variety of inputs.

579 5.2 Debugging Techniques

580 Debugging involves many activities and can be static,
581 dynamic, or post-mortem. *Static debugging* usually
582 takes the form of code review, while *dynamic*
583 *debugging* usually takes the form of tracing and is
584 closely associated with testing. *Post-mortem debugging*
585 is the act of debugging the core dump (memory dump)
586 of a process. Core dumps are often generated after a
587 process has terminated due to an unhandled exception.
588 All three techniques are used at various stages of
589 program development, but (to most programmers)
590 dynamic debugging is the norm.

591 The main activity of dynamic debugging is tracing,
592 which is executing the program one piece at a time,
593 examining the contents of registers and memory, in
594 order to examine the results at each step. There are
595 three ways to trace a program.

- 596 ● *Single-stepping*: execute one instruction at a time
597 to make sure each instruction is executed
598 correctly. This method is tedious but useful in
599 verifying each step of a program.
- 600 ● *Breakpoints*: tell the program to stop executing
601 when it reaches a specific instruction. This
602 technique lets one quickly execute selected code
603 sequences to get a high-level overview of the
604 execution behavior.
- 605 ● *Watch points*: tell the program to stop when a
606 register or memory location changes or when it
607 equals to a specific value. This technique is useful
608 when one doesn't know where or when a value is
609 changed and when this value change likely causes
610 the error.

611 5.3 Debugging Tools

612 Debugging can be complex, difficult, and tedious. Like
613 programming, debugging is also highly creative
614 (sometimes more creative than programming). Thus
615 some help from tools is in order. For dynamic
616 debugging, *debuggers* are widely used and enable the
617 programmer to monitor the execution of a program,
618 stop the execution, re-start the execution, set
619 breakpoints, change values in memory, and even, in
620 some cases, go back in time.

621 For static debugging, there are many so-called *static*
622 *code analysis tools*, which look for a specific set of
623 known problems within the source code. Both
624 commercial and free tools exist in various languages.
625 These tools can be extremely useful when checking
626 very large source trees, where it is impractical to do
627 code walkthroughs. The UNIX *lint* program is an early
628 example.

629 6. Data Structure and Representation 630 [7*s2.1-2.6]

631 Programs work on data. But data must be expressed
632 and organized within computers before being
633 processed by programs. This organization and
634 expression of data for programs' use is the subject
635 of data structure and representation. Simply put, a
636 data structure tries to store and organize data in a
637 computer in such a way that the data can be used
638 efficiently. There are many types of data structures
639 and each type of structure is suitable for some kinds
640 of applications. For example, B/B+ trees are well
641 suited for implementing massive file systems and
642 databases.

643 6.1 Data Structure Overview

644 Data structures are computer representations of data.
645 Data structures are used in almost every program. In a
646 sense, no meaningful program can be constructed
647 without the use of some sort of data structure. Some
648 design methods and programming languages even
649 organize an entire software system around data
650 structures. Fundamentally, data structures are
651 abstractions defined on a collection of data and its
652 associated operations.

653 Often, data structures are designed for improving
654 program or algorithm efficiency. Examples of such
655 data structures include stacks, queues, and heaps. At
656 other times, data structures are used for conceptual
657 unity (abstract data type), such as the name and address
658 of a person. Often, a data structure can determine
659 whether a program runs in a few seconds or in a few
660 hours or even a few days.

661 From the perspective of physical and logical ordering,
662 a data structure is either linear or non-linear. Other
663 perspectives give rise to different classifications that
664 include homogeneous vs. heterogeneous, static vs.
665 dynamic, persistent vs. transit, external vs. internal,
666 primitive vs. aggregate, recursive vs. non-recursive;
667 passive vs. active; and stateful vs. stateless structures.

668 6.2 Types of Data Structure

669 As mentioned above, different perspectives can be used
670 to classify data structures. However, the predominate
671 perspective used in classification centers on physical
672 and logical ordering between data items. This
673 classification divides data structures into linear and
674 non-linear structures. Linear structures organize data
675 items in a single dimension in which each data entry
676 has one (physical or logical) predecessor and one
677 successor with the exception of the first and last entry.
678 The first entry has no predecessor and the last entry has
679 no successor. Non-linear structures organize data items
680 in two or more dimensions, in which case one entry can
681 have multiple predecessors and successors. Examples
682 of linear structures include lists, stacks, and queues.
683 Examples of non-linear structures include heaps, hash
684 tables, and trees (such as binary trees, balance trees, B-
685 trees, and so forth).

686 Another type of data structure that is often encountered
687 in programming is the so-called compound structure. A
688 compound data structure builds on top of other (more
689 primitive) data structures and, in some way, can be
690 viewed as the same structure as the underlying
691 structure. Examples of compound structures include
692 sets, graphs, and partitions. For example, a partition
693 can be viewed as a set of sets.

694 6.3 Operations on Data Structures

695 All data structures support some operations that
696 produce a specific structure and ordering, or
697 retrieve from the structure relevant data, or store
698 data into the structure. Basic operations supported
699 by all data structures include:

- 700 ● Build an initial structure (or initialize the
701 structure).
- 702 ● Insert a data entry into the structure.
- 703 ● Retrieve a data entry from the structure.
- 704 ● Remove a data entry from the structure.

705 Some data structure also support additional
706 operations:

- 707 ● Find a particular element in the structure.

- 708 ● Sort all elements according to some ordering.
- 709 ● Traverse all elements in some specific order.
- 710 ● Reorganize or rebalance the structure.

711 Different structures support different operations
712 with different efficiencies. The difference between
713 operation efficiency can be significant. For example,
714 it is easy to retrieve the last item inserted into a
715 stack, but finding a particular element within a stack
716 is rather slow and tedious.

717 7. Algorithms and Complexity 718 [7* s1.1-1.3, s3.3-3.6, s4.1-4.8, s5.1-5.7, s6.1-6.3, 719 s7.1-7.6, s11.1, s12.1]

720 Programs are not random pieces of code: they are
721 meticulously written to perform user-expected
722 actions. The guide one uses to compose programs
723 are algorithms, which organize various functions
724 into a series of steps and take into consideration the
725 application domain, the solution strategy, and the
726 data structures being used. An algorithm can be
727 very simple or very complex.

728 7.1 Overview of Algorithms

729 Abstractly speaking, algorithms guide the
730 operations of computers and consist of an abstract
731 sequence of actions composed to solve a problem.
732 Alternative definitions include but are not limited
733 to:

- 734 ● An algorithm is any well-defined computational
735 procedure that takes some value or set of values
736 as input and produces some value or set of values
737 as output.
- 738 ● An algorithm is a sequence of computational steps
739 that transform the input into the output.
- 740 ● An algorithm is a tool for solving a well-specified
741 computation problem.

742 Of course, different definitions are favored by
743 different people. Though there is no universally
744 accepted definition, some agreement exists that an
745 algorithm needs to be correct, finite (in other
746 words, terminate eventually), and unambiguous.

747 7.2 Attributes of Algorithms

748 The attributes of algorithms are many and often
749 include modularity, correctness, maintainability,
750 functionality, robustness, user-friendliness,

751 programmer time, simplicity, and extensibility. But
752 the most important attribute is the so-called
753 “performance” or “efficiency,” by which we mean
754 both time and resource-usage efficiency while
755 emphasizing the time axis. To some degree,
756 efficiency determines if an algorithm is feasible or
757 impractical. For example, an algorithm that takes
758 one hundred years to terminate is virtually useless
759 and is even considered incorrect.

760 7.3 Algorithmic Analysis

761 *Analysis of algorithms* is the theoretical study of
762 computer-program performance and resource
763 usage; it determines the goodness of an algorithm.
764 Such analysis usually abstracts away the particular
765 details of a specific computer and focuses on the so-
766 called asymptotic, machine-independent analysis.

767 There are three basic types of analysis. In *worst-*
768 *case analysis*, one determines the maximum time or
769 resources required by the algorithm on any input of
770 size n . In *average-case analysis*, one determines the
771 expected time or resources required by the
772 algorithm over all inputs of size n ; in performing
773 average-case analysis, one often needs to make
774 assumptions on the statistical distribution of inputs.
775 The third type of analysis is the so-called *best-case*
776 *analysis*, in which one determines the minimum
777 time or resources required by the algorithm on any
778 input of size n . Among the three types of analysis,
779 average-case analysis is the most useful but also the
780 most difficult to perform.

781 Besides the basic analysis methods, there are also
782 the *amortized analysis*, in which one determines the
783 maximum time required by an algorithm over a
784 sequence of operations; and the *competitive*
785 *analysis*, in which one determines the relative
786 performance merit of an algorithm against the
787 optimal algorithm (which may not be known) in the
788 same category (for the same operations).

789 7.4 Algorithmic Design Strategies

790 The design of algorithms generally follows one of
791 the following strategies: brute force, divide and
792 conquer, dynamic programming, and greedy
793 selection. The *brute force strategy* is actually a no-
794 strategy. It exhaustively tries every possible way to
795 tackle a problem. If a problem has a solution, this

796 strategy is guaranteed to find it; however, the time
797 expense may be too high. The *divide and conquer*
798 *strategy* improves on the brute force strategy by
799 dividing a big problem into smaller, homogeneous
800 problems. It solves the big problem by recursively
801 solving the smaller problems and combining the
802 solutions to the smaller problems to form the
803 solution to the big problem. The underlying
804 assumption for divide and conquer is that smaller
805 problems are easier to solve.

806 The *dynamic programming strategy* improves on
807 the divide and conquer strategy by recognizing that
808 some of the sub-problems produced by division may
809 be the same and thus avoids solving the same
810 problems again and again. This elimination of
811 redundant sub-problems can dramatically improve
812 efficiency.

813 The *greedy selection strategy* further improves on
814 dynamic programming by recognizing that not all of
815 the sub-problems contribute to the solution of the
816 big problem. By eliminating all but one sub-
817 problem, the greedy selection strategy achieves the
818 highest efficiency among all algorithm design
819 strategies. Sometimes the use of *randomization* can
820 improve on the greedy selection strategy by
821 eliminating the complexity in determining the
822 greedy choice through coin flipping or
823 randomization.

824 7.5 Algorithmic Analysis Strategies

825 The analysis strategies of algorithms include *basic*
826 *counting analysis*, in which one actually counts the
827 number of steps an algorithm takes to complete its
828 task; *asymptotic analysis*, in which one only
829 considers the order of magnitude of the number of
830 steps an algorithm takes to complete its task;
831 *probabilistic analysis*, in which one makes use of
832 probabilities in analyzing the average performance
833 of an algorithm; *amortized analysis*, in which one
834 uses the methods of aggregation, potential, and
835 accounting to analyze the worst performance of an
836 algorithm on a sequence of operations; and
837 *competitive analysis*, in which one uses methods
838 such as potential and accounting to analyze the
839 relative performance of an algorithm to the optimal
840 algorithm.

841 For complex problems and algorithms, one may

842 need to use a combination of the aforementioned
843 analysis strategies.

844 8. Basic Concept of a System 845 [8*c2]

846 Ian Sommerville writes that “a system is a
847 purposeful collection of interrelated components
848 that work together to achieve some objective”
849 [8*p21]. A system can be very simple and include
850 only a few components, like an ink pen, or rather
851 complex, like an aircraft. Depending on whether
852 humans are part of the system, systems can be
853 divided into technical computer-based systems and
854 socio-technical systems, with the former excluding
855 the human factor and the latter including the
856 human factor. Examples of technical, computer-
857 based systems include televisions, mobile phones,
858 and most personal computer software. Examples of
859 socio-technical systems include manned space
860 vehicles, chips embedded inside a human, and so
861 forth.

862 8.1 Emergent System Properties

863 A system is more than simply the sum of its parts.
864 Thus, the properties of a system are not simply the
865 sum of the properties of its components. Instead, a
866 system often exhibits properties that are properties
867 of the system as a whole. These properties are
868 called *emergent properties* because “they emerge
869 only once the system components have been
870 integrated” [8*p23]. Emergent system properties
871 can be either functional or non-functional.
872 Functional properties describe the functions that
873 can be achieved by the system. For example, an
874 aircraft’s functional properties include flotation on
875 air, carrying people or cargo, and use as a weapon
876 of mass destruction. “Non-functional properties
877 relate to the behavior of the system in its
878 operational environment and can include such
879 things as weight, volume, reliability, security,
880 usability, etc.” [8*p23].

881 8.2 System Engineering

882 System engineering is the activity of specifying,
883 designing, implementing, validating, deploying, and
884 maintaining systems. The phases of system
885 engineering vary depending on the system being
886 built but, in general, include requirement definition,

887 system design, sub-system development, system
888 integration, system installation, system evolution,
889 and system decommissioning.

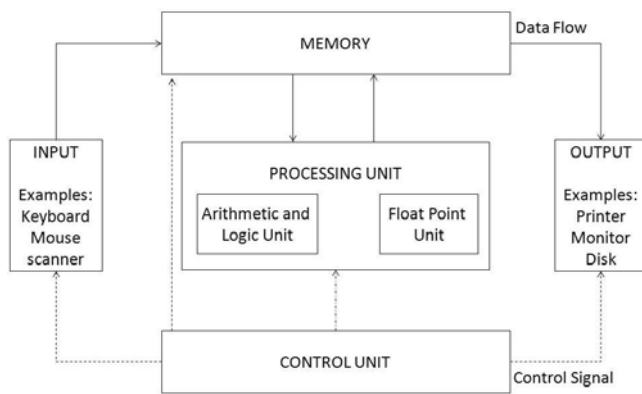
890 Many practical guidelines have been produced in
891 the past to aid people in performing the activities of
892 each phase. For example, system design can be
893 broken into smaller tasks of requirement definition,
894 identification of sub-systems, assignment of
895 requirements to sub-systems, specification of sub-
896 system functionality, definition of sub-system
897 interfaces, and so forth.

898 8.3 Overview of a Computer System

899 Among all the systems, one that is obviously
900 relevant to the software engineering community is
901 the computer system. A computer is a machine that
902 executes programs or software. It consists of a
903 purposeful collection of mechanical, electrical, and
904 electronic components with each component
905 performing a preset function. Jointly, these
906 components are able to perform the computations
907 that are given by a program.

908 Abstractly speaking, a computer receives some
909 input, stores and manipulates some data, and
910 provides some output. The most distinct feature of
911 a computer is its ability to store and execute
912 sequences of instructions called *programs*. An
913 interesting phenomenon concerning the computer
914 is the universal equivalence in functionality.
915 According to Turing, all computers with a certain
916 minimum capability are equivalent in their ability to
917 perform computation tasks. In other words, given
918 enough time and memory, all computers—ranging
919 from a netbook to a supercomputer—are capable of
920 computing exactly the same things, irrespective of
921 speed, size, cost, or anything else.

922 Most computer systems have a structure that is
923 known as the “von Neumann model,” which
924 consists of five components: a *memory* for storing
925 instructions and data, a *processing unit* for
926 performing arithmetic and logical operations, a
927 *control unit* for sequencing and interpreting
928 instructions, *input* for getting external information
929 into the memory, and *output* for producing results
930 for the user. The basic components of a computer
931 system based on the von Neumann model is
932 depicted in Figure 2.



1007 As mentioned before, a computer expresses data
1008 with electrical signals or digital zero and one. Since
1009 there are only two different digits used in data
1010 expression, such a system is called a *binary*
1011 *expression system*. Due to the inherent nature of a
1012 binary system, the maximum numerical value
1013 expressible by an n-bits binary code is $2^n - 1$.
1014 Specifically, binary number $a_n a_{n-1} \dots a_1 a_0$ corresponds
1015 to $a_n \cdot 2^n + a_{n-1} \cdot 2^{n-1} + \dots + a_1 \cdot 2^1 + a_0 \cdot 2^0$. Thus, the numerical
1016 value of the binary expression of 1011 is
1017 $1 \cdot 8 + 0 \cdot 4 + 1 \cdot 2 + 1 \cdot 1 = 11$. To express a non-numerical
1018 value, we need to decide the number of zeros and
1019 ones to use and the order in which those zeros and
1020 ones are arranged.

1021 Of course, there are different ways to do the
1022 encoding, and this gives rise to different data
1023 expression schemes and sub-schemes. For example,
1024 integers can be expressed in the form of unsigned,
1025 one's complement, or two's complement. For
1026 characters, there are ASCII, Unicode, and IBM's
1027 EBCDIC standards. For floating point numbers, there
1028 are IEEE-754 FP 1, 2, and 3 standards.

1029 9.5 The Central Processing Unit (CPU)

1030 The central processing unit is the place where
1031 instructions (or programs) are actually executed.
1032 The execution usually takes several steps, including
1033 fetching the program instruction, decoding the
1034 instruction, fetching operands, performing
1035 arithmetic and logical operations on the operands,
1036 and storing the result. The main components of a
1037 CPU consist of registers where instructions and data
1038 are often read from and written to, the arithmetic
1039 and logic unit (ALU) that performs the actual
1040 arithmetic (such as addition, subtraction,
1041 multiplication, and division) and logic (such as AND,
1042 OR, shift, and so forth) operations, the control unit
1043 that is responsible for producing proper signals to
1044 control the operations, and various (data, address,
1045 and control) buses that link the components
1046 together and transport data to and from these
1047 components.

1048 9.6 Memory System Organization

1049 Memory is the storage unit of a computer. It
1050 concerns the assembling of a large-scale memory
1051 system from smaller and single-digit storage units.
1052 The main topics covered by memory system

1053 architecture include the following:

- 1054 ● Memory cells and chips
- 1055 ● Memory boards and modules
- 1056 ● Memory hierarchy and cache
- 1057 ● Memory as a sub-system of the computer

1058 *Memory cells and chips* deal with single-digit
1059 storage and the assembling of single-digit units into
1060 one-dimensional memory arrays as well as the
1061 assembling of one-dimensional storage arrays into
1062 multi-dimensional storage memory chips. *Memory*
1063 *boards and modules* concern the assembling of
1064 memory chips into memory systems, with the focus
1065 being on the organization, operation, and
1066 management of the individual chips in the system.
1067 *Memory hierarchy and cache* are used to support
1068 efficient memory operations. *Memory as a sub-*
1069 *system* deals with the interface between the
1070 memory system and other parts of the computer.

1071 9.7 Input and Output (I/O)

1072 A computer is useless without I/O. Common input
1073 devices include the keyboard and mouse; common
1074 output devices include the disk, the screen, the
1075 printer, and speakers. Different I/O devices operate
1076 at different data rates and reliabilities. How
1077 computers connect and manage various input and
1078 output devices to facilitate the interaction between
1079 computers and humans (or other computers) is the
1080 focus of topics in I/O. The main issues that must be
1081 resolved in input and output are the ways I/O can
1082 and should be performed.

1083 In general, I/O is performed at both hardware and
1084 software levels. Hardware I/O can be performed in
1085 any of three ways. *Dedicated I/O* dedicates the CPU
1086 to the actual input and output operations during
1087 I/O; *memory-mapped I/O* treats I/O operations as
1088 memory operations; and *hybrid I/O* combines
1089 dedicated I/O and memory-mapped I/O into a single
1090 holistic I/O operation mode.

1091 Coincidentally, software I/O can also be performed
1092 in one of three ways. *Programmed I/O* lets the CPU
1093 wait while the I/O device is doing I/O; *interrupt-*
1094 *driven I/O* lets the CPU's handling of I/O be driven
1095 by the I/O device; and *direct memory access (DMA)*
1096 lets I/O be handled by a secondary CPU embedded
1097 in a so-called DMA device (or channel). (Except

1098 during the initial setup, the main CPU is not
1099 disturbed during a DMA I/O operation.)

1100 Regardless of the types of I/O scheme being used,
1101 the main issues involved in I/O include *I/O*
1102 *addressing* (which deals with the issue of how to
1103 identify the I/O device for a specific I/O operation),
1104 *synchronization* (which deals with the issue of how
1105 to make the CPU and I/O device work in harmony
1106 during I/O), and *error detection and correction*
1107 (which deals with the occurrence of transmission
1108 errors).

1109 10. Compiler Basics 1110 [6*s6.4 9*s8.4]

1111 10.1 Compiler Overview

1112 To be understood by a computer, programs written
1113 in high-level programming languages must be
1114 translated into the specific machine language of the
1115 computer under consideration. This translation is
1116 usually performed by a piece of software called a
1117 compiler. This process of translation from a high-
1118 level language to a machine language is called
1119 compilation (or, sometimes, interpretation).

1120 The primary tasks of a compiler may include
1121 preprocessing, lexical analysis, parsing, semantic
1122 analysis, code generation, and code optimization.
1123 Program faults caused by incorrect compiler
1124 behavior can be very difficult to track down. For this
1125 reason, compiler implementers invest a lot of time
1126 ensuring the correctness of their software.

1127 10.2 Interpretation and Compilation

1128 There are two ways to translate a program written
1129 in a higher-level language into machine code:
1130 interpretation and compilation. *Interpretation*
1131 translates the source code one piece at a time into
1132 machine language, executes it on the spot, and then
1133 goes back for another piece. Both the high-level-
1134 language source code and the interpreter are
1135 required every time the program is run.

1136 *Compilation* translates the high-level-language
1137 source code into an entire machine-language
1138 program (an executable image) by a program called
1139 a compiler. After compilation, only the executable
1140 image is needed to run the program. Most

1141 application software is sold in this form.

1142 10.3 The Compilation Process

1143 Compilation is a complex task. Most compilers
1144 divide the compilation process into many phases. A
1145 typical breakdown is as follows:

- 1146 ● Lexical Analysis
- 1147 ● Syntax Analysis or Parsing
- 1148 ● Semantic Analysis
- 1149 ● Code Generation

1150 *Lexical analysis* partitions the input text (the source
1151 code), which is a sequence of characters, into
1152 separate *comments*, which are to be ignored in
1153 subsequent actions, and *basic symbols*, which have
1154 *lexical meanings*. These basic symbols must
1155 correspond to some terminal symbols of the
1156 grammar of the particular programming language.

1157 *Syntax analysis* is based on the results of the lexical
1158 analysis and discovers the structure in the program
1159 and determines whether or not a text conforms to
1160 an expected format. "Is this a textually correct C++
1161 program?" or "Is this entry textually correct?" are
1162 typical questions that can be answered by syntax
1163 analysis. Syntax analysis determines if the source
1164 code of a program is correct and converts it into a
1165 more structured representation (parse tree) for
1166 semantic analysis or transformation.

1167 *Semantic analysis* adds semantic information to the
1168 parse tree built during the syntax analysis and builds
1169 the symbol table. It performs various semantic
1170 checks that include type checking, object binding
1171 (associating variable and function references with
1172 their definitions), and definite assignment (requiring
1173 all local variables to be initialized before use). If
1174 mistakes are found, the incorrect programs are
1175 rejected.

1176 Once semantic analysis is complete, the phase of
1177 *code generation* begins and transforms the
1178 intermediate code produced in the previous phases
1179 into the native machine language of the computer
1180 under consideration. This involves resource and
1181 storage decisions—such as deciding which variables
1182 to fit into registers and memory and the selection
1183 and scheduling of appropriate machine instructions,
1184 along with their associated addressing modes.

1185 It is often possible to combine multiple phases into
1186 one pass over the code in a compiler
1187 implementation. Some compilers also have a
1188 preprocessing phase at the beginning or after the
1189 lexical analysis that does necessary housekeeping
1190 work, such as processing the program instructions
1191 for the compiler (directives). Some compilers
1192 provide an optional optimization phase at the end
1193 of the entire compilation to optimize the code (such
1194 as the rearrangement of instruction sequence) for
1195 efficiency and other desirable objectives requested
1196 by the users.

1197 **11. Operating Systems Basics** 1198 **[6*c3]**

1199 Every system of meaningful complexity needs to be
1200 managed. A computer, as a rather complex
1201 electrical-mechanical system, needs its own
1202 manager for managing the resources and activities
1203 occurring on it. That manager is called an *operating*
1204 *system* (OS).

1205 *11.1 Operating Systems Overview*

1206 Operating systems manage and beautify computers
1207 such that they can be correctly and easily used by
1208 humans. Conceptually, an operating system is a
1209 computer program that manages the hardware
1210 resources and makes it easier to use by applications
1211 by presenting nice abstractions. This nice
1212 abstraction is often called the virtual machine and
1213 includes such things as processes, virtual memory,
1214 and file systems. An OS hides the complexity of the
1215 underlying hardware and is found on almost any
1216 modern computer.

1217 The principal roles played by OSs are management
1218 and illusion. *Management* refers to the OS's
1219 management (allocation and recovery) of physical
1220 resources among multiple competing
1221 users/applications/tasks. *Illusion* refers to the nice
1222 abstractions the OS provides.

1223 *11.2 Tasks of an Operating System*

1224 The tasks of an operating system differ significantly
1225 depending on the machine and time of its invention.
1226 However, modern operating systems have come to
1227 agreement as to the tasks that must be performed by an
1228 OS. These tasks include CPU management, memory
1229 management, disk management (file system), I/O

1230 device management, and security and protection. Each
1231 OS task manages one type of physical resource.

1232 Specifically, CPU management deals with the
1233 allocation and releases of the CPU among competing
1234 programs (called processes/threads in OS jargon),
1235 including the operating system itself. The main
1236 abstraction provided by CPU management is the
1237 process/thread model. Memory management deals with
1238 the allocation and release of memory space among
1239 competing processes, and the main abstraction
1240 provided by memory management is virtual memory.
1241 Disk management deals with the sharing of magnetic
1242 disks among multiple programs/users and its main
1243 abstraction is file system. I/O device management deals
1244 with the allocation and releases of various I/O devices
1245 among competing processes. Security and protection
1246 deal with the protection of computer resources from
1247 illegal use.

1248 *11.3 Operating System Abstractions*

1249 The arsenal of OSs is abstraction. Corresponding to the
1250 five physical tasks, OSs use five abstractions:
1251 process/thread, virtual memory, file systems,
1252 input/output, and protection domains. The overall OS
1253 abstraction is the virtual machine.

1254 For each task area of OS, there is both a physical
1255 reality and a conceptual abstraction. The physical
1256 reality refers to the hardware resource under
1257 management; the conceptual abstraction refers to
1258 the interface the OS presents to the users/programs
1259 above. For example, in the thread model of the OS,
1260 the physical reality is the CPU and the abstraction is
1261 multiple CPUs. Thus, a user doesn't have to worry
1262 about sharing the CPU with others when working on
1263 the abstraction provided by an OS. In the virtual
1264 memory abstraction of an OS, the physical reality is
1265 the physical RAM or ROM (whatever), the
1266 abstraction is multiple unlimited memory space.
1267 Thus, a user doesn't have to worry about sharing
1268 physical memory with others or about limited
1269 physical memory size.

1270 *11.4 Operating Systems Classification*

1271 Different operating systems can have different
1272 functionality implementation. In the early days of
1273 the computer era, operating systems were relatively
1274 simple. As time goes on, the complexity and
1275 sophistication of operating systems increases
1276 significantly. From a historical perspective, operating

1277 system can be classified as one of the following.

- 1278 ● *Batching OS*: organizes and processes work in
1279 batches. Examples of such OSs include IBM's
1280 FMS, IBSYS, and University of Michigan's
1281 UMES.
- 1282 ● *Multi-programmed batching OS*: adds multi-task
1283 capability into earlier simple batching OSs. An
1284 example of such an OS is IBM's OS/360.
- 1285 ● *Time-sharing OS*: adds multi-task and interactive
1286 capabilities into the OS. Examples of such OSs
1287 include UNIX, Linux, and NT.
- 1288 ● *Real-time OS*: adds timing predictability into the
1289 OS by scheduling individual tasks according to
1290 each task's completion deadlines. Examples of
1291 such OS include VxWorks (WindRiver) and
1292 DART (EMC).
- 1293 ● *Distributed OS*: adds the capability of managing a
1294 network of computers into the OS.
- 1295 ● *Embedded OS*: has limited functionality and is
1296 used for embedded systems such as cars and
1297 PDAs. Examples of such OSs include Palm OS,
1298 Windows CE, and TOPPER.

1299 Alternatively, an OS can be classified by its
1300 applicable target machine/environment into the
1301 following.

- 1302 ● *Mainframe OS*: runs on the mainframe computers
1303 and include OS/360, OS/390, AS/400, MVS, and
1304 VM.
- 1305 ● *Server OS*: runs on workstations or servers and
1306 includes such systems as UNIX, Windows, Linux,
1307 and VMS.
- 1308 ● *Multi-computer OS*: runs on multiple computers
1309 and include such examples as Novell Netware.
- 1310 ● *Personal computers OS*: runs on personal
1311 computers and include such examples as DOS,
1312 Windows, Mac OS, and Linux.

1313 The most popular modern operating systems today
1314 include Microsoft Windows, UNIX, Linux, and Mac
1315 OS X.

1316 **12. Database Basics and Data Management** 1317 **[6*c9]**

1318 A database consists of an organized collection of data
1319 for one or more uses. In a sense, a database is a
1320 generalization and expansion of data structures. But the
1321 difference is that a database is usually external to
1322 individual programs and permanent in existence
1323 compared to data structures. Databases are used when
1324 the data volume is large or logical relations between
1325 data items are important. The factors considered in

1326 database design include performance, concurrency,
1327 integrity, and recovery from hardware failures.

1328 **12.1 Entity and Schema**

1329 The things a database tries to model and store are
1330 called entities. Entities can be real-world objects
1331 such as persons, cars, houses, and so forth,. or
1332 abstract concepts such as human, salary, names,
1333 and so forth. An entity can be primitive such as a
1334 name or composite (such as an employee that
1335 consists of a name, identification number, salary,
1336 address, and so forth).

1337 The single most important concept in a database is
1338 the so-called *schema*, which is a description of the
1339 entire database structure from which all other
1340 database activities are built. A schema defines the
1341 relationships between the various entities that
1342 compose a database. For example, a schema for a
1343 company payroll system would consist of such
1344 things as employee ID, name, salary rate, address,
1345 and so forth. Database software maintains the
1346 database according to the schema.

1347 Another important concept in database is the so-called
1348 *database models* that describe the type of relationship
1349 among various entities. The commonly used models
1350 include relational, network, and object models.

1351 **12.2 Database Management Systems (DBMS)**

1352 The software that manages the database is called
1353 the database management system (DBMS). A DBMS
1354 controls the creation, maintenance, and use of the
1355 database and is usually categorized according to the
1356 database model it supports—such as the relational,
1357 network, or object model. For example, a relational
1358 database management system (RDBMS) implements
1359 features of the relational model. An object database
1360 management system (ODBMS) implements features
1361 of the object model.

1362 **12.3 Database Query Language**

1363 Users/applications interact with a database through
1364 a database query language, which is a specialized
1365 programming language tailored to database use.
1366 The database model tends to determine the query
1367 languages that are available to access the database.
1368 One commonly used query language for the

relational database is the structured query language, more commonly abbreviated as SQL. A common query language for object databases is the object query language (abbreviated as OQL), although not all vendors implement this. An example of an SQL query may look like:

```
SELECT Component_No, Quantity
FROM COMPONENT
WHERE Item_No = 100
```

The above query selects all the Component_No and its corresponding quantity from a database table called COMPONENT, where the Item_No equals to 100.

12.4 Tasks of DBMS Packages

A DBMS system provides the following capabilities:

- *Database development* is used to define and organize the content, relationships, and structure of the data needed to build a database.
- *Database interrogation* is used for accessing the data in a database for information retrieval and report generation. End users can selectively retrieve and display information and produce printed reports. This is the operation that most users know about databases.
- *Database Maintenance* is used to add, delete, update, correct, and protect the data in a database.
- *Application Development* is used to develop prototypes of data entry screens, queries, forms, reports, tables, and labels for a prototyped application. It also refers to the use of 4GL (or 4th Generation Language) or application generators to develop or generate program code.

12.5 Data Management

A database must manage the data stored in it. This management includes both organization and storage.

The organization of the actual data in a database depends on the database model. In a relational model, data are organized as tables with different tables representing different entities or relations among a set of entities. The storage of data deals with the storage of these database tables on disks. The common ways for achieving this is to use files. Sequential, indexed, and hash files are all used in this purpose with different file structures providing different access performance and convenience.

12.6 Data Mining

One often has to know what to look for before querying a database. This type of access does not make full use of the vast amount of information stored in the database, and in fact reduces the database into a collection of discrete records that are not related with each other. To take full advantage of a database, one can perform statistical analysis and pattern discovery on the content of a database using a technique called *data mining*. Such operations can be used to support a number of business activities that include, but are not limited to, marketing, fraud detection, and trend analysis.

Numerous ways for performing data mining have been invented in the past decade and include such common techniques as class description, class discrimination, cluster analysis, association analysis, and outlier analysis.

13. Network Communication Basics [9*c12]

A computer network connects a collection of computers and allows users of different computers to share resources with other users. A network facilitates the communications between all the connected computers and may give the illusion of a single, omnipotent computer. Every computer or device connected to a network is called a *network node*.

A number of computing paradigms have emerged to benefit from the functions and capabilities provided by computer networks. These paradigms include distributed computing, grid computing, Internet computing, and cloud computing.

13.1 Types of Network

Computer networks are not all the same and may be classified according to a wide variety of characteristics, including the network's connection method, wired technologies, wireless technologies, scale, network topology, functions, and speed. But the classification that is familiar to most is based on the scale of networking.

- *Personal Area Network/Home Network* is a computer network used for communication

among computer(s) and different information technological devices close to one person. The devices connected to such a network may include PCs, faxes, PDAs, and TVs. This is the base on which the Internet of Things is built.

- *Local Area Network* (LAN) connects computers and devices in a limited geographical area, such as a school campus, computer laboratory, office building, or closely positioned group of buildings.
- *Campus Network* is a computer network made up of an interconnection of local area networks (LANs) within a limited geographical area.
- *Wide area network* (WAN) is a computer network that covers a large geographic area, such as a city or country or even across intercontinental distances. A WAN limited to a city is sometimes called a Metropolitan Area Network.
- *Internet* is the global network that connects computers located in many (perhaps all) countries.

Other classifications may divide networks into control networks, storage networks, virtual private networks (VPN), wireless networks, point-to-point networks, and Internet of Things.

13.2 Basic Network Components

All networks are made up of the same basic hardware components, including computers, network interface cards (NICs), bridges, hubs, switches, and routers. All these components are called *nodes* in the jargon of networking. Each component performs a distinctive function that is essential for the packaging, connection, transmission, amplification, controlling, unpacking, and interpretation of the data. For example, a repeater amplifies the signals, a switch performs many-to-many connections, a hub performs one-to-many connections, an interface card is attached to the computer and performs data packing and transmission, a bridge connects one network with another, and a router is a computer itself and performs data analysis and flow control to regulate the data from the network.

The functions performed by various network components correspond to the functions specified by one or more levels of the seven-layer OSI networking model depicted in Figure 4.

13.3 Networking Protocols and Standards

Computers talk with each other using protocols, which specify the format and regulations used to pack and un-pack data. To facilitate easier communication and better structure, network

protocols are divided into different layers with each layer dealing with one aspect of the communication. For example, the physical layers deal with the physical connection between the parties that are to communicate, the data link layer deals with the raw data transmission and flow control, and the network layer deals with the packing and un-packing of data into a particular format that is understandable by the relevant parties. The most commonly used OSI networking model organizes network protocols into seven layers, as depicted in Figure 4.

Figure 4. The seven-layer OSI networking model

Application Layer
Presentation Layer
Session Layer
Transport Layer
Network Layer
Data link Layer
Physical Layer

One thing to note is that not all network protocols implement all layers of the OSI model. For example, the TCP/IP protocol implements neither the presentation layer nor the session layer.

There can be more than one protocol for each layer. For example, UDP and IP both work on the network layer. Physical layer protocols include token ring, Ethernet, fast Ethernet, gigabit Ethernet, and wireless. Data link layer protocols include frame-relay, asynchronous transfer mode (ATM), and Point-to-Point Protocol (PPP). Application layer protocols include Fibre channel, Small Computer System Interface (SCSI), and Bluetooth. For each layer or even each individual protocol, there may be standards established by national or international organizations to guide the design and development of the corresponding protocols.

13.4 The Internet

The Internet is a global system of interconnected governmental, academic, corporate, public, and private computer networks. Most people gain access to the Internet through the service of so-called Internet service providers (ISP). The ISP maintains one or more switching centers called a point of presence, which is often known to users through their domain names such as .com, .gov, .org, .net, and .cn.

1547 13.5 *Internet of Things*

1548 The Internet of Things refers to the networking of
1549 everyday objects—such as cars, cell phones, PDAs,
1550 TVs, refrigerators, and even buildings—using wired
1551 or wireless networking technologies. The function
1552 and purpose of *Internet of Things* is to interconnect
1553 all things to facilitate autonomous and better living.
1554 Technologies used in the Internet of Things include
1555 RFID, wireless and wired networking, sensor
1556 technology, and much software of course. As the
1557 paradigm of Internet of Things is still shaping up,
1558 many works are needed for Internet of Things to
1559 actually gain wide spread acceptance.

1560 13.6 Virtual Private Network (VPN)

1561 A *Virtual Private Network* is a computer network in
1562 which some or all of the links between the nodes
1563 are carried by open connections or virtual circuits in
1564 some larger network (for example, the Internet).
1565 VPNs are most often used to separate the traffic of
1566 different user communities and often provide better
1567 security to users.

1568 14. Parallel and Distributed Computing 1569 [9*c9]

1570 Parallel computing is a computing paradigm that
1571 emerges with the development of multi-functional units
1572 within a computer. The main objective of parallel
1573 computing is to execute several tasks simultaneously
1574 on different functional units and thus improve
1575 throughput or response or both. With the emergence of
1576 computer networks, parallel computing has also taken
1577 on a new perspective in which the parallelism occurs
1578 on a greater scale. Distributed computing, on the other
1579 hand, is a computing paradigm that emerges with the
1580 development of computer networks. Its main objective
1581 is to either make use of multiple computers in the
1582 network to accomplish things otherwise not possible
1583 within a single computer or improve computation
1584 efficiency by harnessing the power of multiple
1585 computers.

1586 14.1 *Parallel and Distributed Computing Overview*

1587 Traditionally, parallel computing investigates ways to
1588 maximize concurrency (the simultaneous execution of
1589 multiple tasks) within the boundary of a computer.
1590 Distributed computing studies distributed systems,
1591 which consists of multiple *autonomous* computers that
1592 communicate through a computer network.
1593 Alternatively, distributed computing can also refer to

1594 the use of distributed systems to solve computational or
1595 transactional problems. In the former definition,
1596 distributed computing investigates the protocols,
1597 mechanisms, and strategies that provide the foundation
1598 for distributed computation; in the latter definition,
1599 distributed computing studies the ways of dividing a
1600 problem into many tasks and assigning such tasks to
1601 various computers involved in the computation.

1602 Fundamentally, distributed computing is another
1603 form of parallel computing, albeit on a grander
1604 scale. In distributed computing, the functional units
1605 are not ALU, FPU, or separate cores, but individual
1606 computers. For this reason, some people regard
1607 distributed computing as being the same as parallel
1608 computing. Because both distributed and parallel
1609 computing involve some form of concurrency, they
1610 are both also called concurrent computing.

1611 14.2 *Difference between Parallel and Distributed* 1612 *Computing*

1613 Though parallel and distributed computing
1614 resemble each other on the surface, there is a
1615 subtle but real distinction between them: parallel
1616 computing does not necessarily refer to the
1617 execution of programs on different computers—
1618 instead, they can be run on different processors
1619 within a single computer. In fact, consensus among
1620 computing professionals limits the scope of parallel
1621 computing to the case where a shared memory is
1622 used by all processors involved in the computing,
1623 while distributed computing refers to computations
1624 where private memory exists for each processor
1625 involved in the computations.

1626 Thus, it is possible to classify concurrent systems as
1627 being "parallel" or "distributed" based on the
1628 existence or non-existence of shared memory
1629 among all the processors. Thus, parallel computing
1630 deals with computations within a single computer;
1631 distributed computing deals with computations
1632 within a set of computers. According to this view,
1633 multicore computing is a form of parallel
1634 computing.

1635 14.3 *Parallel and Distributed Computing Models*

1636 Since multiple computers/processors/cores are
1637 involved in distributed/parallel computing, some
1638 coordination among the involved parties is
1639 necessary to ensure correct behavior of the system.

1640 Different ways of coordination give rise to different
1641 computing models. The most common models in
1642 this regard are the shared memory (parallel) model
1643 and the message-passing (distributed) model.

1644 In a *shared memory (parallel)* model, all computers
1645 have access to a shared memory. The algorithm
1646 designer chooses the program for execution by each
1647 computer. Access to the memory can be
1648 synchronous or asynchronous, and must be
1649 coordinated such that coherency is maintained.
1650 Different access models have been invented for
1651 such a purpose.

1652 In a *message-passing (distributed)* model, all
1653 computers run the same program. The system must
1654 work correctly regardless of the structure of the
1655 network. This model can be further classified into
1656 client-server (CS), browser-server (BS), and many-
1657 layers (ML) models. In the CS model, the server
1658 provides services and the client requests services
1659 from the server. In the so-called BS model, the
1660 server provides services and the client is the
1661 browser. In the ML model, each layer provides
1662 services to the layer immediately above it and
1663 requests services from the layer immediately below
1664 it. In fact, the ML model can be seen as a chain of
1665 client-server models. Often, the layers between the
1666 bottommost layer and the topmost layer are called
1667 *middleware*, which is a distinct subject of study in
1668 its own right.

1669 14.4 Main Issues in Distributed Computing

1670 Coordination among all the components in a
1671 distributed computing environment is often
1672 complex and time-consuming. As the number of
1673 cores/CPU's/computers increases, the complexity of
1674 distributed computing also increases. Among the
1675 many issues faced, memory coherency and
1676 consensus among all computers are the most
1677 difficult ones. Many computation paradigms have
1678 been invented to solve these problems and are the
1679 main discussion issues in distributed/parallel
1680 computing.

1681 15. Basic User Human Factors 1682 [4*c8, 5*c5]

1683 Software is developed to meet human desires or
1684 needs. Thus, all software design and development

1685 must take into consideration human-user factors
1686 such as how people use software, how people view
1687 software, and what humans expect from software.
1688 There are numerous factors in the human-machine
1689 interaction, but the **basic** human-user factors
1690 considered here include input/output, the handling
1691 of error messages, and the robustness of the
1692 software in general.

1693 15.1 Input and Output

1694 Input and output are the interfaces between users
1695 and software. Software is useless without input and
1696 output. Humans design software to process some
1697 input and produce desirable output. All software
1698 engineers must consider input and output as an
1699 integral part of the software product they engineer
1700 or develop. Issues considered for input include (but
1701 are not limited to):

- 1702 ● What input is required?
- 1703 ● How is the input passed from users to computers?
- 1704 ● What is the most convenient way for users to
1705 enter input?
- 1706 ● What format does the computer require of the
1707 input data?

1708 For output, we need to consider what the users
1709 wish to see:

- 1710 ● In what format would users like to see output?
- 1711 ● What is the most pleasing way to display output?

1712 Furthermore, if the party receiving the output is not
1713 human but another computer or control system,
1714 then we need to consider the output type and
1715 format that the software should produce to ensure
1716 proper data feed into another system.

1717 There are many rules of thumb for developers to
1718 follow to produce good input/output for a system.
1719 These rules of thumb include simple and natural
1720 dialogue, speaking users' language, minimizing user
1721 memory load, and consistency.

1722 15.2 Error Messages

1723 It is understandable that software contains bugs and
1724 fails from time to time. But users should be notified
1725 if there is anything that impedes the smooth
1726 execution of the program. Nothing is more

frustrating than an unexpected termination or behavioral deviation of software without any warning or explanation. To be user friendly, the software should report all error conditions to the users or upper-level applications so that some measure can be taken to rectify the situation or to exit gracefully. There are several guidelines that define what constitutes a good error message: error messages should be clear, to the point, and timely.

First, error messages should clearly explain what is happening so that users know what is going on in the system. Second, error messages should pinpoint the cause of the error, if at all possible, so that proper actions can be taken. Third, error messages should be displayed right when the error condition occurs. According to Jakob Nielsen, “Good error messages should be expressed in plain language (no codes), precisely indicate the problem, and constructively suggest a solution” [5*p20].

15.3 Software Robustness

Software robustness refers to the ability of software to tolerate erroneous inputs. Software is said to be robust if it continues to function even when erroneous inputs are given. Thus, it is unacceptable for software to simply crash when encountering an input problem as this breaks the abstraction and may cause unexpected consequences, such as the loss of valuable data. Software that exhibits such behavior is considered to lack robustness.

Nielsen gives a simpler (though not very accurate) description of software robustness: “The software should have a low error rate, so that users make few errors during the use of the system and so that if they do make errors they can easily recover from them. Further, catastrophic errors must not occur” [5*p26].

There are many ways to evaluate the robustness of software and just as many ways to make software more robust. For example, to improve robustness, one should always check the validity of the inputs and return values before progressing further; one should always throw an exception when something unexpected occurs, and one should never quit a program without first giving users/applications a chance to correct the condition..

16. Basic Developer Human Factors [4*c31-32]

Developer human factors refer to the considerations of human factors taken when developing software. Software is developed by humans, read by humans, and maintained by humans. If anything is wrong, humans are responsible for correcting those wrongs. Thus, it is essential to write software in a way that is easily understandable by humans or, at the very least, by other software developers. A program that is easy to read and understand exhibits a so-called readability.

The means to ensure that software meet this objective are numerous and range from proper architecture at the macro level to the particular coding style and variable usage at the micro level. But the two prominent factors are *structure* (or program layouts) and *comments* (documentation).

16.1 Structure

Most people prefer structure to chaos. It is human nature to prefer a well-structured program to randomly organized sections of code. Structure not only looks good but also makes programs easier to understand and maintain. It is said that a well-structured program is self-explanatory. But if a program is poorly structured, then no amount of explanation or comments is sufficient to make it understandable. The ways to organize a program are numerous and range from the proper use of white space, indentation, and parentheses to nice arrangements of groupings, blank lines, and braces. Whatever style one chooses, it should be consistent across the entire program.

16.2 Comments

To most people, programming is coding. These people do not realize that programming also includes writing comments and that comments are an integral part of programming. True, comments are not used by the computer and certainly do not constitute final instructions for the computer, but they improve the readability of the programs by explaining the meaning and logic of the statements or sections of code. It should be remembered that programs are not only meant for computers, they are also read, written, and modified by humans.

The types of comments include repeat of the code, explanation of the code, marker of the code, summary of the code, description of the code's intent, and information that cannot possibly be expressed by the code itself. The best comments are self-documenting code. If the code is written in such a clear and precise manner that its meaning is self-proclaimed, then no comment is needed. But this is easier said than done. Most programs are not self-explanatory and are often hard to read and understand if no comments are given.

Here are some general guidelines for writing good comments:

- Comments should be consistent across the entire software.
- Each function should be associated with comments that explain the purpose of the function and its role in the overall software.
- Within a function, comments should be given for each logical section of coding to explain the meaning and purpose (intention) of the section.
- Comments are seldom required for individual statements. If a statement needs comments, one should reconsider the statement.

17. Secure Coding [10*c29]

Due to increasing malicious activities targeted at computer systems, security has become a significant issue in the development of software systems. In addition to the usual correctness and reliability, software developers must also pay attention to the security of the software they develop. Secure coding is one of the ways to ensure security.

17.1 Two Aspects of Secure Coding

Secure coding means different things for different people. It can mean the way a specific function is coded, such that the coding itself is secure, or it can mean the coding of security into software systems.

Most people entangle the two together without distinction. One reason for such entanglement is that it is not clear how one can make sure that a specific coding is secure. For example, in C programming language, the expression of $i < 1$ (shift the value of variable i to the left by 1 bit) and $2 * i$ (multiply the value of variable i by constant 2) mean the same thing semantically, but do they have the

same security ramification? Due to this lack of understanding, secure coding—in its current state of existence—mostly refers to the second aspect mentioned above: the coding of security into software system.

17.2 Coding Security into Software

A generally accepted view concerning software security is that it is much better to design security into software than to patch it in after software is developed. To design security into software, one must take into consideration most, if not every, stage of the software development life cycle. In particular, secure coding involves *requirement security*, *design security*, and *implementation security*, which are described in the next few sections.

17.3 Requirement Security

Requirement security deals with the clarification and specification of security policy and objectives into software requirements, which lays the foundation for all future security consideration in the software development process. Factors to consider in this phase include requirements and threats. The former refers to the specific functions that are required for the sake of security; the latter refers to the possible ways that the security of software is threatened.

17.4 Design Security

Design security deals with the construction of software modules that fit together to meet the security objective specified in the security requirements. This step clarifies the details of security considerations and develops the specific steps for implementation. Factors considered may include frameworks and access modes that set up the overall security monitoring/enforcement strategies, as well as the individual policy enforcement mechanisms.

17.5 Implementation Security

Implementation security is directly related to coding and is the most relevant to secure coding among the three phases of coding security into software. All secure requirements and designs are useless if

1906 the implementation is not secure. Implementation
1907 security concerns itself with the question of how to
1908 write actual codes for specific situations such that
1909 security considerations are taken care of.

1956
1957

1958

Science. Boston: Addison-Wesley, 2002.

1910 Implementation security can be achieved by
1911 following some recommended rules. A few such
1912 rules follow [10*pp880-910]:

- 1913 ● Structure the process so that all sections requiring
1914 extra privileges are modules. The modules should
1915 be as small as possible and should perform only
1916 those tasks that require those privileges.
- 1917 ● Ensure that any assumptions in the program are
1918 validated. If this is not possible, document them
1919 for the installers and maintainers so they know the
1920 assumptions that attackers will try to invalidate.
- 1921 ● Ensure that the program does not share objects in
1922 memory with any other program.
- 1923 ● The error status of every function must be
1924 checked. Do not try to recover unless neither the
1925 cause of the error nor its effects affect any
1926 security considerations. The program should
1927 restore the state of the system to the state it had
1928 before the process began, and then terminate.

- 1929 [1] The Joint Task Force on Computing
1930 Curricula, *et al.*, "Software Engineering
1931 2004: Curriculum Guidelines for
1932 Undergraduate Degree Programs in
1933 Software Engineering," 2004.
- 1934 [2] S. T. Frezza, "Computer Science: Is It Really
1935 the Scientific Foundation for Software
1936 Engineering," *Computer*, vol. 43, no. 8 (Aug.
1937 2010), pp. 98-101.
- 1938 [3*] G. Voland, *Engineering by Design*, 2nd ed.
1939 Upper Saddle River, NJ: Prentice Hall, 2003.
- 1940 [4*] S. McConnell, *Code Complete*, 2nd ed.
1941 Redmond, WA: Microsoft Press, 2004.
- 1942 [5*] J. Nielsen, *Usability Engineering*, 1st ed.
1943 Boston: Morgan Kaufmann, 1993.
- 1944 [6*] J. G. Brookshear, *Computer Science: An
1945 Overview*, 10th ed. Boston: Addison-
1946 Wesley, 2008.
- 1947 [7*] E. Horowitz, *et al.*, *Computer Algorithms*,
1948 2nd ed. Summit, NJ: Silicon Press, 2007.
- 1949 [8*] I. Sommerville, *Software Engineering*, 8th
1950 ed. New York: Addison-Wesley, 2006.
- 1951 [9*] L. Null and J. Lobur, *The Essentials of
1952 Computer Organization and Architecture*,
1953 2nd ed. Sudbury, MA: Jones and Bartlett
1954 Publishers, 2006.
- 1955 [10*] M. Bishop, *Computer Security: Art and*

1959
1960

Matrix of Topics vs. References for the Computing Foundations KA

	[1*]	[2*]	[3*]	[4*]	[5*]	[6*]	[7*]	[8*]
1. Problem Solving Techniques								c3, s4.1-4.2 -- 35
1.1 Definition of Problem Solving								pp
1.2 Formulating the Real Problem								s3.1
1.3 Analyze the Problem								s3.2
1.4 Design a Solution Search Strategy								s3.3
1.5 Problem Solving using Programs					c5			s4.2
2. Abstraction					s5.2-5.4 -- 41			
2.1 Levels of Abstraction					pp			
2.2 Encapsulation					s5.2-5.3			
2.3 Hierarchy					s5.3			
3. Programming Fundamentals					s5.2			
3.1 The Programming Process					c6-19 -- 338			
3.2 Programming Paradigms					pp			
3.3 Defensive Programming					c6-c19			
4. Programming Language Basics					c6-c19			
4.1 Programming Language Overview		c6 -- 60			c8			
4.2 Syntax and Semantics of Programming Language		pp						
4.3 Low Level Programming Language		s6.1						
4.4 High Level Programing Language		S6.2						
4.5 Declarative vs. Imperative Programming Language								
5. Debugging Tools and Techniques								
5.1 Types of Errors					c23 -- 28 pp			
5.2 Debugging Techniques:					s23.1			
5.3 Debugging Tools					s23.2			
6. Data Structure and Representation					s23.5			
6.1 Data Structure Overview			s2.1-2.6 -- 47 pp					
6.2 Types of Data Structure			s2.1-2.6					
6.3 Operations on Data Structures			s2.1-2.6					
7. Algorithms and Complexity			s1.1-1.3, s3.3-3.6, s4.1-4.8, s5.1-5.7, s6.1-6.3, s7.1-7.6, s11.1, s12.1 --					

7.1 Overview of Algorithms	254 pp	
7.2 Attributes of Algorithms	s1.1-1.2	
7.3 Algorithmic Analysis	s1.3	
7.4 Algorithmic Design Strategies	s1.3	
	s3.3-3.6, s4.1-4.8, s5.1-5.7, s6.1-6.3, s7.1-7.6, s11.1, s12.1	
7.5 Algorithmic Analysis Strategies	s3.3-3.6, s4.1-4.8, s5.1-5.7, s6.1-6.3, s7.1-7.6, s11.1, s12.1	
8. Basic Concept of A System		c2-- 21
8.1 Emergent System Properties		pp
8.2 System Engineering		s2.1
8.3 Overview of a Computer System		s2.2
9. Computer Organization		c1-4 -- 96
9.1 Computer Organization Overview		pp
9.2 Digital Systems		s1.1-1.2
9.3 Digital Logic		c3
9.3 Computer Expression of Data		c3
9.4 The Central Processing Unit (CPU)		c2
9.5 Memory System Organization		s4.1-4.2
9.6 Input and Output (I/O)		s4.6
10. Compiler Basics	s6.4	s4.5
10.1 Compiler Overview		s8.4 -- 11 pp
10.2 Interpretation and Compilation		s8.4
10.3 The Compilation Process	s6.4	s8.4-8.5
11. Operating Systems Basics	c3 -- 26	s8.4
	pp	
11.1 Operating Systems Overview	s3.2	
11.2 Tasks of Operating System	s3.3	
11.3 Operating System Abstractions	s3.2	
11.4 Operating Systems Classification	s3.1	
12. Database Basics and Data Management	c9 -- 44	
	pp	
12.1 Entity and Schema	s9.1	
12.2 Database Management Systems (DBMS)	s9.1	
12.3 Database Query Language	s9.2	
12.4 Tasks of DBMS Packages	s9.2	

12.5 Data Management	s9.5		
12.6 Data Mining	s9.6		
13. Network Communication Basics			c12 -- 70 pp
13.1 Types of Network			s12.2-12.3
13.2 Basic Network Components			s12.6
13.3 Networking Protocols and Standards			s12.4-12.5
13.4 The Internet			
13.5 Internet of Things			s12.8
13.6 Virtual Private Network			
14. Parallel and Distributed Computing			c9 -- 65 pp
14.1 Parallel and Distributed Computing Overview			s9.4.1-9.4.3
14.2 Difference between Parallel and Distributed Computing			s9.4.4-9.4.5
14.3 Parallel and Distributed Computing Models			s9.4.4-9.4.5
14.4 Main Issues in Distributed Comp.			
15. Basic User Human Factors		c8 -- 28 pp	c5 -- 50 pp
15.1 Input and Output			s5.1, s5.3
15.2 Error Messages			s5.2, s5.8
15.3 Software Robustness			s5.5-5.6
16. Basic Developer Human Factors		c31-32 -- 90	
16.1 Structure		pp	
16.2 Comments		c31	
		c32	
17. Secure Coding	c29 -- 55		
	pp		
17.1 Two Aspects of Secure Coding	s29.1		
17.2 Coding Security into Software	s29.4		
17.3 Requirement Security	s29.2		
17.4 Design Security	s29.3		
17.5 Implementation Security	s29.5		
Total Pages Referenced in the Computing Foundations KA			1367 pages
1961			
1962	Recommended references for the Knowledge Area		
1963	[1*] M. Bishop, <i>Computer Security: Art and Science</i> . Boston: Addison-Wesley, 2002.		
1964	[2*] J. G. Brookshear, <i>Computer Science: An Overview</i> , 10th ed. Boston: Addison-Wesley, 2008.		
1965	[3*] E. Horowitz, et al., <i>Computer Algorithms</i> , 2nd ed. Summit, NJ: Silicon Press, 2007.		
1966	[4*] S. McConnell, <i>Code Complete</i> , 2nd ed. Redmond, WA: Microsoft Press, 2004.		

- 1967 [5*] J. Nielsen, *Usability Engineering*, 1st ed. Boston: Morgan Kaufmann, 1993.
- 1968 [6*] L. Null and J. Lobur, *The Essentials of Computer Organization and Architecture*, 2nd ed. Sudbury, MA: Jones and Bartlett Publishers, 2006.
- 1969 [7*] I. Sommerville, *Software Engineering*, 8th ed. New York: Addison-Wesley, 2006.
- 1970 [8*] G. Voland, *Engineering by Design*, 2nd ed. Upper Saddle River, NJ: Prentice Hall, 2003.
- 1971

1972 Further Readings List for the Computing Foundations KA [1–7]

- 1973
- 1974 1. A.V. Aho, R. Sethi, and J. Ullman, *Compilers: Principles, Techniques, and Tools*, Addison Wesley, 2006.
- 1975 Reason for inclusion: The book is considered by many as one of the primary sources for knowledge on
- 1976 compiler and its construction. It provides more in-depth discussion on compiler, in general, and the
- 1977 compilation process, in particular.
- 1978
- 1979 2. T.H. Cormen, et al., *Introduction to Algorithm*, 3rd ed., The MIT Press, 2009.
- 1980 Reason for inclusion: The book is considered by many as one of the primary sources for knowledge on
- 1981 algorithms and is used by many universities as either a textbook or a reference book. It is brought in to
- 1982 provide an alternative view or perspective on algorithms.
- 1983
- 1984 3. ACM SIGPLAN Education Board, "Why Undergraduates Should Learn the Principles of Programming
- 1985 Languages," 2010.
- 1986 Reason for inclusion: This paper is cited in the KA description.
- 1987
- 1988 4. S.T. Frezza, "Computer Science: Is It Really the Scientific Foundation for Software Engineering?" *Computer*,
- 1989 IEEE CS Press, vol. 43, no. 8, Aug. 2010, pp. 98–101.
- 1990 Reason for inclusion: This paper is cited in the KA description.
- 1991
- 1992 5. Graduate Software Engineering 2009, "Curriculum Guidelines for Graduate Degree Programs in Software
- 1993 Engineering," Aug. 2010, www.gswe2009.org.
- 1994 Reason for inclusion: These guidelines are cited in the KA description.
- 1995
- 1996 6. A. Silberschatz, P.B. Galvin, and G. Gagne, *Operating System Concepts*, 8th ed., Wiley, 2008.
- 1997 Reason for inclusion: This is a commonly used textbook in many American universities in an operating systems
- 1998 course. It provides more in-depth discussion on operating systems basics and is already included in the
- 1999 consolidated reference list for KAs other than the Computing Foundations KA.
- 2000
- 2001 7. The Joint Task Force on Computing Curricula, IEEE Computer Society, and Association for Computing
- 2002 Machinery, "Software Engineering 2004: Curriculum Guidelines for Undergraduate Degree Programs in
- 2003 Software Engineering," 2004.
- 2004 Reason for inclusion: These guidelines are cited in the KA description.
- 2005